

# Project2

Group-11

```
library(tidyverse)
library(moderndiver)
library(gapminder)
library(sjPlot)
library(jtools)
library(GGally)
library(gt)
library(gridExtra)
library(knitr)
library(patchwork)
library(broom)
library(MASS)
library(janitor)
library(pscl)
library(ggfortify)
```

## 1 Data Wrangling

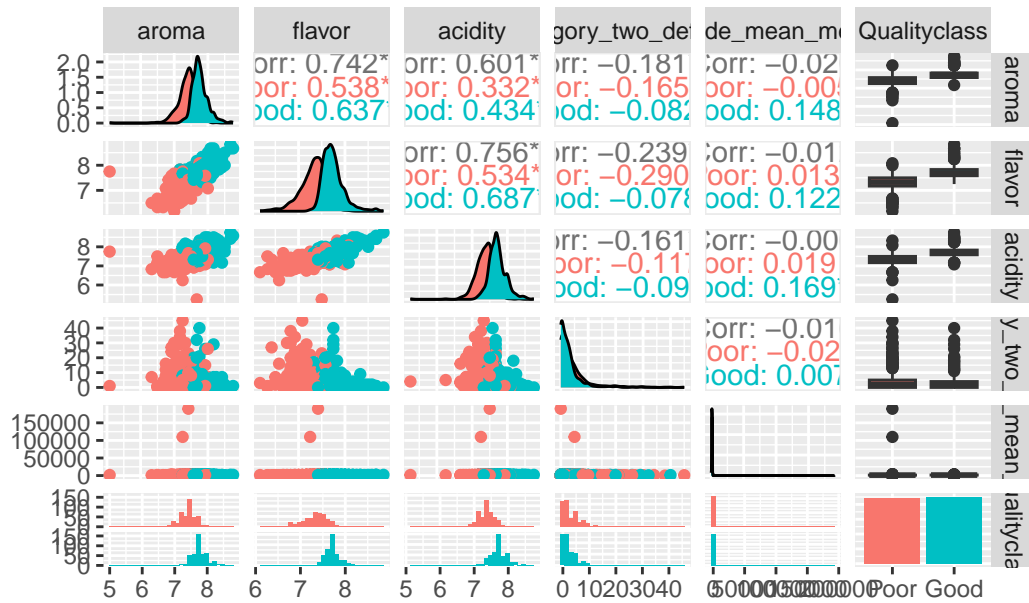
Preprocess the data and conduct summary statistics.

```
# Read the dataset
Data <- read.csv("dataset11.csv")
Data <- na.omit(Data)
Data$Qualityclass <- factor(Data$Qualityclass, levels = c("Poor", "Good"))
Data$harvested <- factor(Data$harvested, levels = 2010:2018)

# Scatterplot matrix with ggpairs()
scatterplot = Data %>%
  dplyr::select(aroma, flavor, acidity, category_two_defects, altitude_mean_meters, Quality)
```

```
ggpairs(scatterplot, aes(color = Qualityclass), title="Scatterplot matrix with ggpairs()")
```

Scatterplot matrix with ggpairs()



```
# Remove outliers
q1_aroma <- quantile(Data$aroma, 0.25)
q3_aroma <- quantile(Data$aroma, 0.75)
iqr_aroma <- q3_aroma - q1_aroma
lower_bound_aroma <- q1_aroma - 1.5 * iqr_aroma
upper_bound_aroma <- q3_aroma + 1.5 * iqr_aroma
Data1 <- Data %>%
  filter(aroma >= lower_bound_aroma & aroma <= upper_bound_aroma)

q1_flavor <- quantile(Data1$flavor, 0.25)
q3_flavor <- quantile(Data1$flavor, 0.75)
iqr_flavor <- q3_flavor - q1_flavor
lower_bound_flavor <- q1_flavor - 1.5 * iqr_flavor
upper_bound_flavor <- q3_flavor + 1.5 * iqr_flavor
Data1 <- Data1 %>%
  filter(flavor >= lower_bound_flavor & flavor <= upper_bound_flavor)

q1_acidity <- quantile(Data1$acidity, 0.25)
q3_acidity <- quantile(Data1$acidity, 0.75)
iqr_acidity <- q3_acidity - q1_acidity
```

```

lower_bound_acidity <- q1_acidity - 1.5 * iqr_acidity
upper_bound_acidity <- q3_acidity + 1.5 * iqr_acidity
Data1 <- Data1 %>%
  filter(acidity >= lower_bound_acidity & acidity <= upper_bound_acidity)

q1_defects <- quantile(Data1$category_two_defects, 0.25)
q3_defects <- quantile(Data1$category_two_defects, 0.75)
iqr_defects <- q3_defects - q1_defects
lower_bound_defects <- q1_defects - 1.5 * iqr_defects
upper_bound_defects <- q3_defects + 1.5 * iqr_defects
Data1 <- Data1 %>%
  filter(category_two_defects >= lower_bound_defects & category_two_defects <= upper_bound_defects)

q1_altitude <- quantile(Data1$altitude_mean_meters, 0.25)
q3_altitude <- quantile(Data1$altitude_mean_meters, 0.75)
iqr_altitude <- q3_altitude - q1_altitude
lower_bound_altitude <- q1_altitude - 1.5 * iqr_altitude
upper_bound_altitude <- q3_altitude + 1.5 * iqr_altitude
Data1 <- Data1 %>%
  filter(altitude_mean_meters >= lower_bound_altitude & altitude_mean_meters <= upper_bound_altitude)

# Standardize the 'altitude_mean_meters' column
mean_altitude <- mean(Data1$altitude_mean_meters)
sd_altitude <- sd(Data1$altitude_mean_meters)
Data1$altitude_mean_meters <- (Data1$altitude_mean_meters - mean_altitude) / sd_altitude

```

## 2 Data Visualization

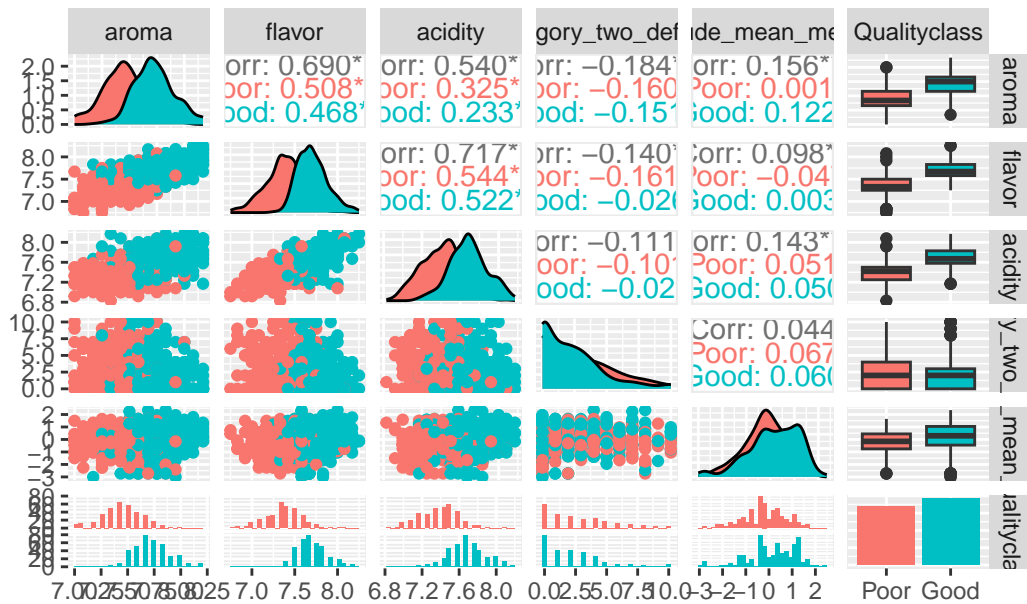
Generate visualizations to better understand the data.

```

# ggpairs of the wrangling data
scatterplot = Data1 %>%
  dplyr::select(aroma, flavor, acidity, category_two_defects, altitude_mean_meters, Qualityclass)
ggpairs(scatterplot, aes(color = Qualityclass), title="Scatterplot matrix with ggpairs()")

```

## Scatterplot matrix with ggpairs()



```
# Summary Statistics for 'aroma' and 'flavor' across different quality classes
Data1 |>
```

```
  summarize('ar.Mean' = mean(aroma),
            'ar.Sd' = sd(aroma),
            'ar.Min' = min(aroma),
            'ar.Max' = max(aroma),
            'fl.Mean' = mean(flavor),
            'fl.Sd' = sd(flavor),
            'fl.Min' = min(flavor),
            'fl.Max' = max(flavor),
            .by = Qualityclass) |>
```

```
gt() |>
```

```
  fmt_number(decimals = 2) |>
```

```
  tab_spanner(
```

```
    label = "aroma",
```

```
    columns = c(ar.Mean, ar.Sd, ar.Min, ar.Max)
```

```
) |>
```

```
  tab_spanner(
```

```
    label = "flavor",
```

```
    columns = c(fl.Mean, fl.Sd, fl.Min, fl.Max)
```

```
)
```

```
# Summary statistics for 'acidity' and 'category_two_defects' across different quality classes
Data1 |>
```

```

summarize('ac.Mean' = mean(acidity),
          'ac.Sd' = sd(acidity),
          'ac.Min' = min(acidity),
          'ac.Max' = max(acidity),
          'C.Mean' = mean(category_two_defects),
          'C.Sd' = sd(category_two_defects),
          'C.Min' = min(category_two_defects),
          'C.Max' = max(category_two_defects),
          .by = Qualityclass) |>
gt() |>
  fmt_number(decimals = 2) |>
  tab_spanner(
    label = "acidity",
    columns = c(ac.Mean, ac.Sd, ac.Min, ac.Max)
  ) |>
  tab_spanner(
    label = "Defects",
    columns = c(C.Mean, C.Sd, C.Min, C.Max)
  )
# Summary statistics for 'altitude_mean_meters' across different quality classes
Data1 |>
  summarize('A.Mean' = mean(altitude_mean_meters),
            'A.Sd' = sd(altitude_mean_meters),
            'A.Min' = min(altitude_mean_meters),
            'A.Max' = max(altitude_mean_meters),
            .by = Qualityclass) |>
gt() |>
  fmt_number(decimals = 2) |>
  tab_spanner(
    label = "Altitude mean meters",
    columns = c(A.Mean, A.Sd, A.Min, A.Max)
  )

# Calculate the count of coffee bean qualities for each country
quality_counts <- Data1 %>%
  group_by(country_of_origin, Qualityclass) %>%
  summarise(count = n()) %>%
  spread(Qualityclass, count, fill = 0) %>%
  mutate(proportion_good = Good / (Good + Poor))

# Create a bar plot showing the proportion of good quality coffee beans by country

```

Table 1: Summary statistics

(a)

Qualityclass	aroma				flavor			
	ar.Mean	ar.Sd	ar.Min	ar.Max	fl.Mean	fl.Sd	fl.Min	fl.Max
Poor	7.44	0.19	7.00	8.00	7.36	0.21	6.75	8.08
Good	7.73	0.18	7.17	8.17	7.71	0.17	7.25	8.25


(b)

Qualityclass	acidity				Defects			
	ac.Mean	ac.Sd	ac.Min	ac.Max	C.Mean	C.Sd	C.Min	C.Max
Poor	7.38	0.20	6.83	8.08	2.75	2.64	0.00	10.00
Good	7.69	0.20	7.17	8.17	2.25	2.37	0.00	10.00

(c)

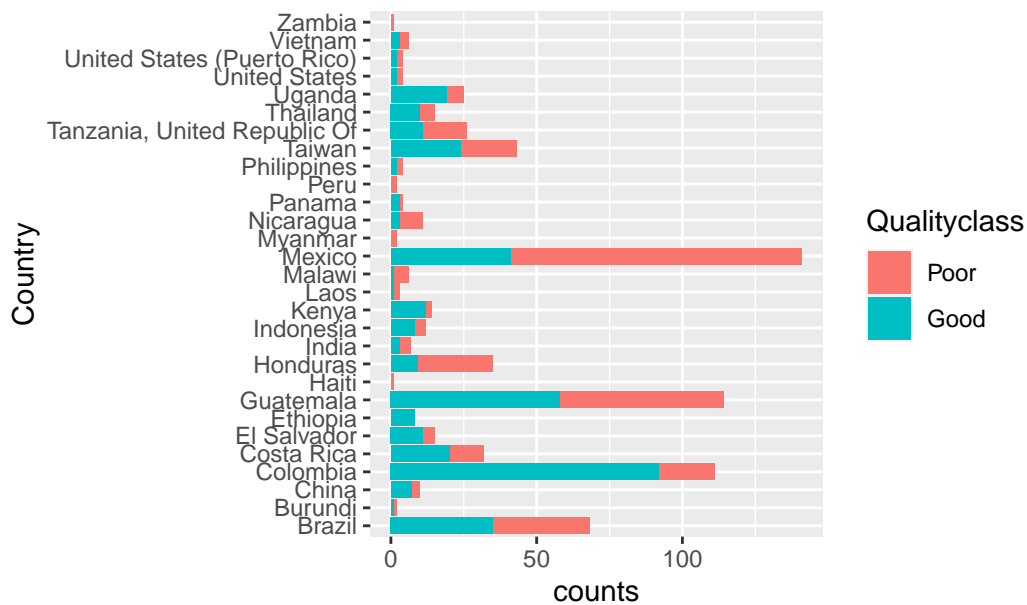
Qualityclass	Altitude mean meters			
	A.Mean	A.Sd	A.Min	A.Max
Poor	−0.18	0.91	−2.73	1.65
Good	0.16	1.05	−3.00	2.35

```
ggplot(quality_counts, aes(x = country_of_origin, y = proportion_good, fill = country_of_o
  geom_bar(stat = "identity", show.legend = FALSE) +
  labs(x = "Country", y = "Proportion of Good Quality Coffee Beans",
       title = "Proportion of Good Quality Coffee Beans by Country") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust = 1))
```




```
# Create a bar plot to visualize the distribution of coffee bean quality by country
counts <- Data1 %>%
  group_by(country_of_origin, Qualityclass) %>%
  summarise(count = n())
ggplot(data = counts, mapping = aes(x = country_of_origin, y = count, fill = Qualityclass))
  geom_col() +
  labs(x = "Country", y = "counts",
       title = "Distribution of coffee bean quality by country")+
  coord_flip()
```

Distribution of coffee bean quality by country

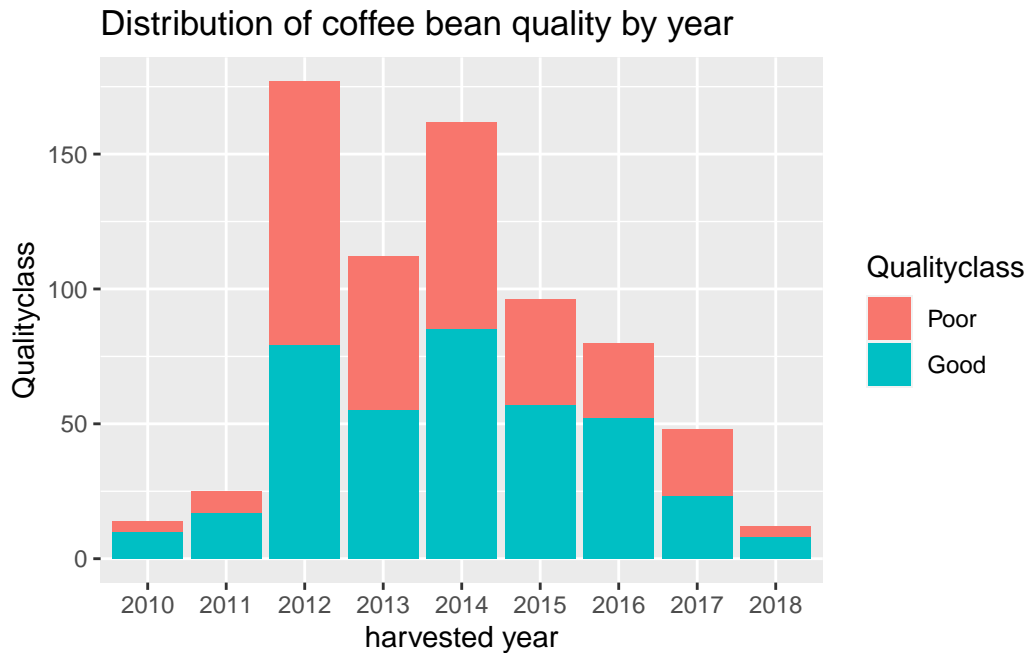


```
# Create a bar plot showing the proportion of good quality coffee beans by year
quality_counts1 <- Data1 %>%
  group_by(harvested, Qualityclass) %>%
  summarise(count = n()) %>%
  spread(Qualityclass, count, fill = 0) %>%
  mutate(proportion_good = Good / (Good + Poor))
ggplot(quality_counts1, aes(x =harvested, y = proportion_good, fill = harvested)) +
  geom_bar(stat = "identity",show.legend = FALSE) +
  labs(x = "harvested year", y = "Proportion of Good Quality Coffee Beans",
       title = "Proportion of Good Quality Coffee Beans by year") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=1))
```





```
# Create a bar plot to visualize the distribution of coffee bean quality by year
counts1 <- Data1 %>%
  group_by(harvested, Qualityclass) %>%
  summarise(count = n())
ggplot(data = counts1, mapping = aes(x = harvested, y = count, fill = Qualityclass)) +
  geom_col() +
  labs(x = "harvested year", y = "Qualityclass",
       title = "Distribution of coffee bean quality by year")
```



### 3 Exploratory Data Analysis

```
# Conduct an origin model
model_full <- glm(Qualityclass ~ country_of_origin + aroma + flavor + acidity + category_t
                  family = binomial(link = "logit"))
model_full %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ country_of_origin + aroma + flavor +
     acidity + category_two_defects + altitude_mean_meters + harvested,
     family = binomial(link = "logit"), data = Data1)
```

Coefficients:

	Estimate	Std. Error	z value
(Intercept)	-155.93383	13.71901	-11.366
country_of_originBurundi	1.92335	5.32186	0.361
country_of_originChina	0.51607	1.23498	0.418
country_of_originColombia	1.79012	0.63394	2.824
country_of_originCosta Rica	0.38187	0.87151	0.438

country_of_originEl Salvador	0.17324	0.97069	0.178
country_of_originEthiopia	12.19028	1178.76284	0.010
country_of_originGuatemala	-0.82776	0.61013	-1.357
country_of_originHaiti	-13.36267	3956.18039	-0.003
country_of_originHonduras	-1.10992	0.80864	-1.373
country_of_originIndia	-3.07110	1.13658	-2.702
country_of_originIndonesia	-0.68124	1.19806	-0.569
country_of_originKenya	0.02514	1.77726	0.014
country_of_originLaos	1.13360	1.96356	0.577
country_of_originMalawi	-0.65762	1.41622	-0.464
country_of_originMexico	-0.89996	0.57619	-1.562
country_of_originMyanmar	-14.36618	2797.39843	-0.005
country_of_originNicaragua	0.05486	1.82141	0.030
country_of_originPanama	3.38904	1.78721	1.896
country_of_originPeru	-18.66679	2192.69041	-0.009
country_of_originPhilippines	2.80520	3.13355	0.895
country_of_originTaiwan	0.62404	0.78671	0.793
country_of_originTanzania, United Republic Of	0.93548	0.91419	1.023
country_of_originThailand	2.18907	0.95049	2.303
country_of_originUganda	-1.48521	0.86689	-1.713
country_of_originUnited States	1.84776	2.01063	0.919
country_of_originUnited States (Puerto Rico)	-1.41603	1.49613	-0.946
country_of_originVietnam	1.67138	1.29784	1.288
country_of_originZambia	-13.01366	3956.18042	-0.003
aroma	6.03872	0.99701	6.057
flavor	8.29375	1.15796	7.162
acidity	6.21276	0.98315	6.319
category_two_defects	0.11822	0.05970	1.980
altitude_mean_meters	0.24303	0.18054	1.346
harvested2011	-0.39748	1.21849	-0.326
harvested2012	0.01028	1.06170	0.010
harvested2013	0.41447	1.06015	0.391
harvested2014	0.52414	1.08744	0.482
harvested2015	0.43049	1.07645	0.400
harvested2016	1.34805	1.14006	1.182
harvested2017	1.29547	1.13936	1.137
harvested2018	2.35384	1.52305	1.545
Pr(> z )			
(Intercept)	< 2e-16 ***		
country_of_originBurundi	0.71780		
country_of_originChina	0.67603		
country_of_originColombia	0.00475 **		
country_of_originCosta Rica	0.66126		

country_of_originEl Salvador	0.85835
country_of_originEthiopia	0.99175
country_of_originGuatemala	0.17487
country_of_originHaiti	0.99731
country_of_originHonduras	0.16988
country_of_originIndia	0.00689 **
country_of_originIndonesia	0.56961
country_of_originKenya	0.98871
country_of_originLaos	0.56372
country_of_originMalawi	0.64240
country_of_originMexico	0.11831
country_of_originMyanmar	0.99590
country_of_originNicaragua	0.97597
country_of_originPanama	0.05792 .
country_of_originPeru	0.99321
country_of_originPhilippines	0.37067
country_of_originTaiwan	0.42765
country_of_originTanzania, United Republic Of	0.30617
country_of_originThailand	0.02127 *
country_of_originUganda	0.08666 .
country_of_originUnited States	0.35810
country_of_originUnited States (Puerto Rico)	0.34391
country_of_originVietnam	0.19781
country_of_originZambia	0.99738
aroma	1.39e-09 ***
flavor	7.93e-13 ***
acidity	2.63e-10 ***
category_two_defects	0.04767 *
altitude_mean_meters	0.17827
harvested2011	0.74427
harvested2012	0.99228
harvested2013	0.69583
harvested2014	0.62981
harvested2015	0.68922
harvested2016	0.23703
harvested2017	0.25553
harvested2018	0.12223

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.53 on 725 degrees of freedom

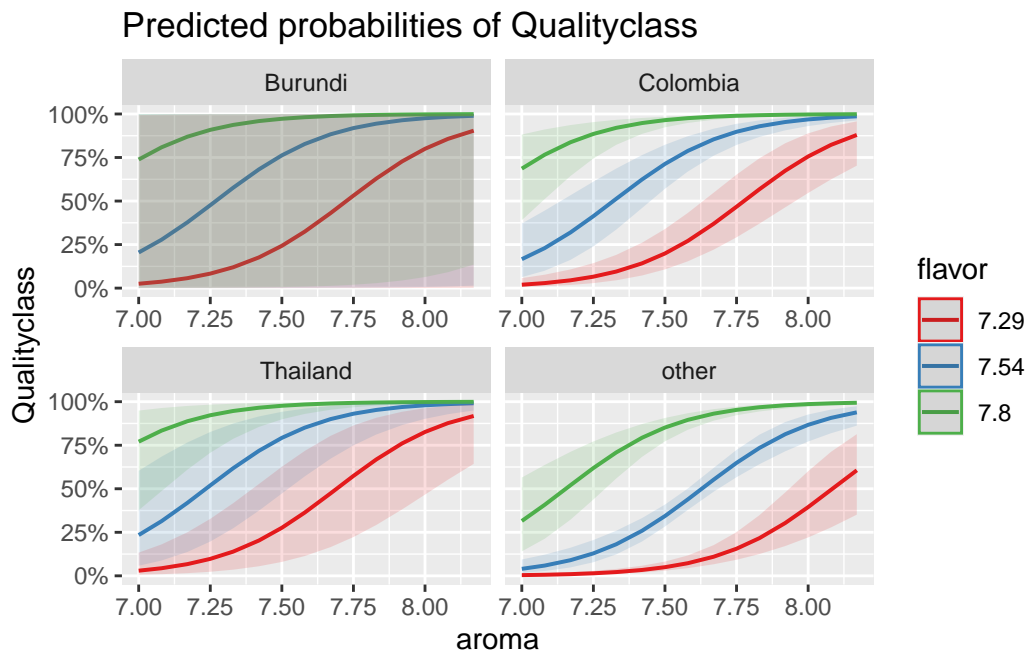
Residual deviance: 361.12 on 684 degrees of freedom  
AIC: 445.12

Number of Fisher Scoring iterations: 16

```
# Create a new variable 'country_category' based on 'country_of_origin' and convert 'country_of_origin' to factor
data <- Data1 %>%
  mutate(country_category = ifelse(country_of_origin %in% c("Burundi", "Colombia", "Thailand"), country_of_origin, "other"))
data$country_category <- factor(data$country_category, levels = c("Burundi", "Colombia", "Thailand", "other"))
```

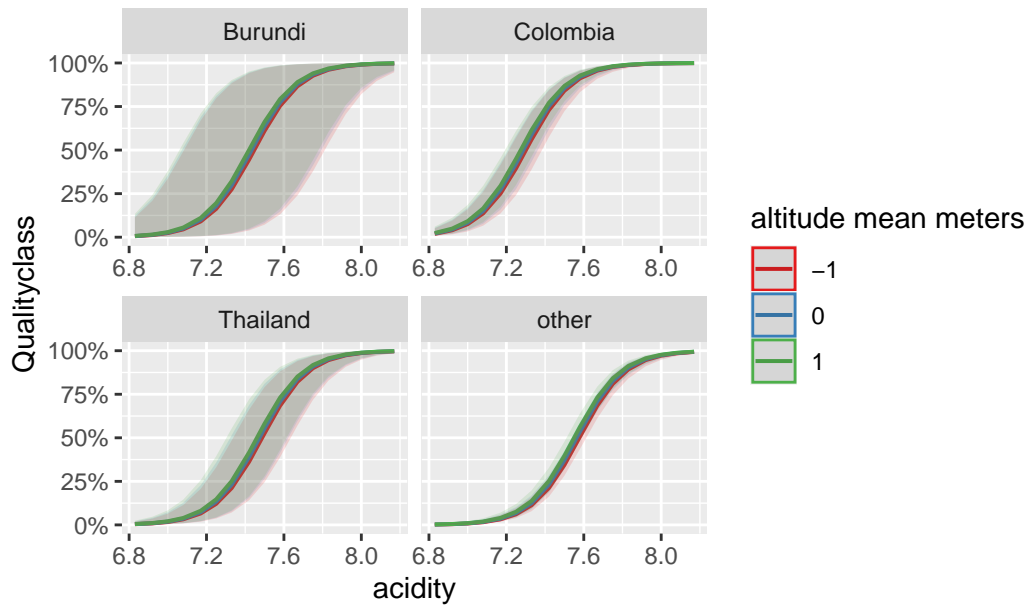
### 3.1 Interaction effects plots

```
model_country_aroma_flavor <- glm(Qualityclass ~ country_category + aroma + flavor, data = data)
plot_model(model_country_aroma_flavor, type = "pred", terms = c("aroma", "flavor", "country_category"))
```



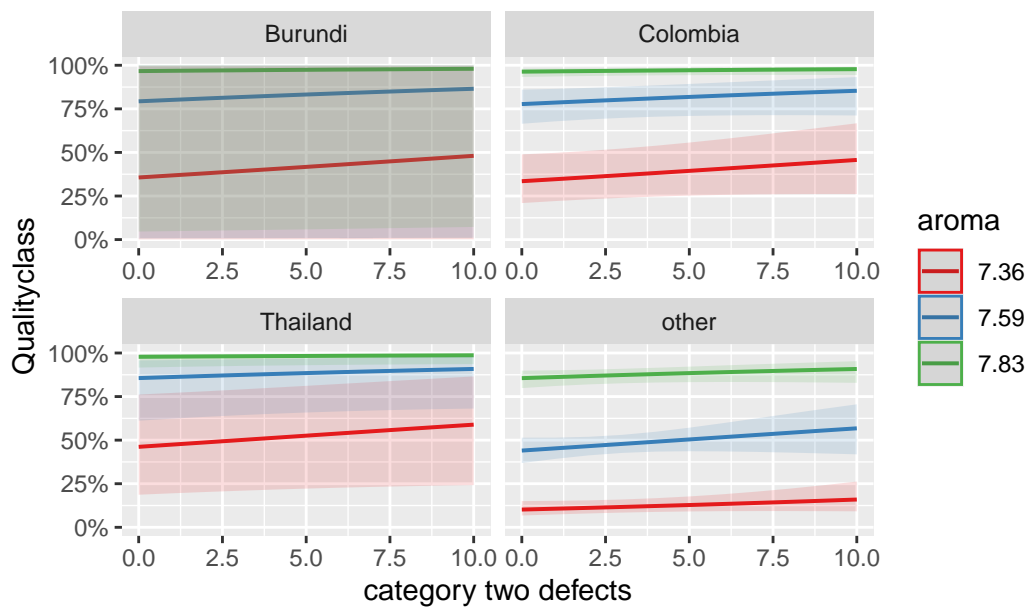
```
model_country_acidity_altitude <- glm(Qualityclass ~ country_category + acidity + altitude, data = data)
plot_model(model_country_acidity_altitude, type = "pred", terms = c("acidity", "altitude_mean_meters", "country_category"))
```

Predicted probabilities of Qualityclass



```
model_aroma_flavor_acidity <- glm(Qualityclass ~ country_category + aroma + category_two_d
plot_model(model_aroma_flavor_acidity, type = "pred",
           terms = c("category_two_defects", "aroma", "country_category"))
```

Predicted probabilities of Qualityclass



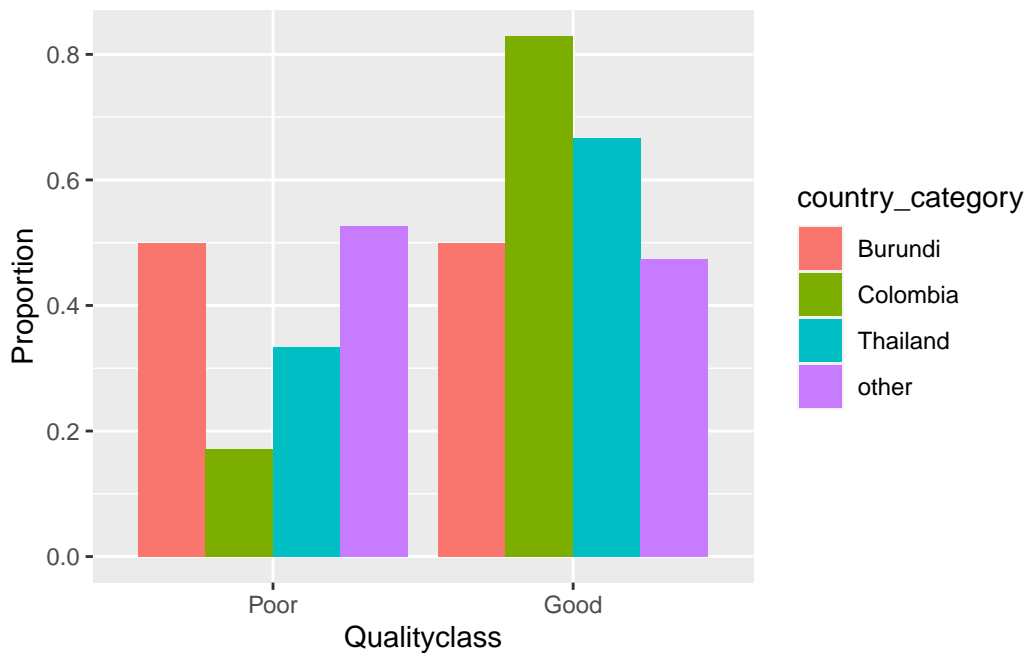
Modeling each predictor separately with the response variable to observe the individual impact of each feature on the quality of coffee.

## 3.2 Country and Qualityclass

```
# Select 'country_category' and 'Qualityclass' columns and generate a contingency table.
data_country_category <- data %>%
  dplyr::select(country_category, Qualityclass)
data_country_category %>%
  tabyl(country_category, Qualityclass) %>%
  adorn_percentages() %>%
  adorn_pct_formatting() %>%
  adorn_ns()
```

country_category		Poor		Good	
Burundi	50.0%	(1)	50.0%	(1)	
Colombia	17.1%	(19)	82.9%	(92)	
Thailand	33.3%	(5)	66.7%	(10)	
other	52.7%	(315)	47.3		(283)

```
# Create a barplot of 'country_category' across different 'Qualityclass' levels
p0 <- ggplot(data_country_category, aes(x = Qualityclass, y = after_stat(prop), group = country_category)) +
  geom_bar(position = "dodge", stat = "count") +
  labs(y = "Proportion")
p0
```



```
# Fit logistic regression model with 'country_category' predictor and 'Qualityclass' response
model_country <- glm(Qualityclass ~ country_category, data = data_country_category, family = binomial)
summary(model_country)
```

Call:

```
glm(formula = Qualityclass ~ country_category, family = binomial(link = "logit"),
     data = data_country_category)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-8.713e-14	1.414e+00	0.000	1.000
country_categoryColombia	1.577e+00	1.436e+00	1.098	0.272
country_categoryThailand	6.931e-01	1.517e+00	0.457	0.648
country_categoryother	-1.071e-01	1.417e+00	-0.076	0.940

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.53 on 725 degrees of freedom  
 Residual deviance: 950.78 on 722 degrees of freedom  
 AIC: 958.78



Number of Fisher Scoring iterations: 4

```
# Extract coefficients from the model and calculate their confidence intervals.
model_country_coef_logodds <- model_country %>%
  summary() %>%
  coef()
model_country_coef_logodds
```

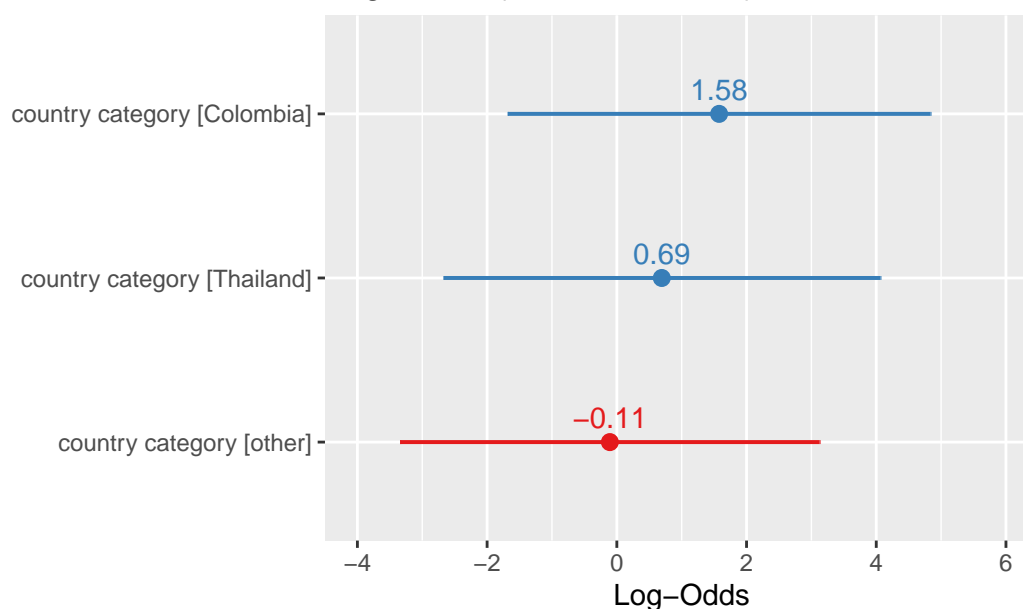
	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-8.712937e-14	1.414214	-6.160977e-14	1.0000000
country_categoryColombia	1.577350e+00	1.436489	1.098059e+00	0.2721788
country_categoryThailand	6.931472e-01	1.516575	4.570477e-01	0.6476367
country_categoryother	-1.071257e-01	1.416583	-7.562262e-02	0.9397193

```
confint_logodds <- confint(model_country)
confint_logodds
```

	2.5 %	97.5 %
(Intercept)	-3.230337	3.230337
country_categoryColombia	-1.681791	4.837248
country_categoryThailand	-2.670586	4.065181
country_categoryother	-3.340551	3.126299

```
# Plot log-odds of being a good instructor
plot_model(model_country, show.values = TRUE, transform = NULL,
           title = "Log-Odds (Good instructor)", show.p = FALSE)
```

## Log-Odds (Good instructor)




```
# Transform the coefficients into odds ratios and obtain their confidence intervals
model_country_coef_odds <- model_country %>%
  summary() %>%
  coef() %>%
  exp()
model_country_coef_odds
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	1.0000000	4.113250	1.000000	2.718282
country_categoryColombia	4.8421053	4.205904	2.998340	1.312822
country_categoryThailand	2.0000000	4.556592	1.579404	1.911019
country_categoryother	0.8984127	4.123009	0.927166	2.559263

```
exp(confint_logodds)
```


	2.5 %	97.5 %
(Intercept)	0.03954417	25.28818
country_categoryColombia	0.18604041	126.12174
country_categoryThailand	0.06921165	58.27545
country_categoryother	0.03541742	22.78949

```
# Plot odds of being a good instructor
plot_model(model_country, show.values = TRUE,
           title = "Odds (Good instructor)", show.p = FALSE)
```



```
# Calculate log odds, odds, and probabilities and store them
data_country_category_after <- data_country_category %>%
  mutate(logodds.Good = predict(model_country, type = "response")) %>%
  mutate(odds.Good = exp(logodds.Good)) %>%
  mutate(probs.Good = fitted(model_country))


# Generate a predictive plot
plot_model(model_country, type = "pred",
           terms = c("country_category"))
```



### 3.3 Aroma and Qualityclass

```
# Select 'aroma' and 'Qualityclass' columns
data_aroma <- data %>%
  dplyr::select(aroma, Qualityclass)

# Create a boxplot of 'aroma' across different 'Qualityclass' levels
p1 <- ggplot(data = data_aroma, aes(x = Qualityclass, y = aroma, fill = Qualityclass)) +
  geom_boxplot() +
  labs(x = "Qualityclass", y = "aroma")+
  theme(legend.position = "none")
p1
```



```
# Fit logistic regression model with 'aroma' predictor and 'Qualityclass' response
model1 <- glm(Qualityclass ~ aroma, data = data_aroma,
              family = binomial(link = "logit"))
model1 %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ aroma, family = binomial(link = "logit"),
    data = data_aroma)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-63.1944	4.8098	-13.14	<2e-16 ***
aroma	8.3465	0.6342	13.16	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.53 on 725 degrees of freedom  
Residual deviance: 676.23 on 724 degrees of freedom

AIC: 680.23


Number of Fisher Scoring iterations: 5

```
# Calculate lower and upper bounds for 'aroma' log-odds
mod1.coef.logodds <- model1 %>%
  summary() %>%
  coef()
aroma.logodds.lower <- mod1.coef.logodds["aroma", "Estimate"] -
  1.96 * mod1.coef.logodds["aroma", "Std. Error"]
aroma.logodds.upper <- mod1.coef.logodds["aroma", "Estimate"] +
  1.96 * mod1.coef.logodds["aroma", "Std. Error"]

# Display the confidence interval
paste("(", aroma.logodds.lower, ",", aroma.logodds.upper, ")")
```

```
[1] "( 7.1035545966061 , 9.58952426960609 )"
```

```
# Plot log-odds of being a good instructor
plot_model(model1, show.values = TRUE, transform = NULL,
  title = "Log-Odds (Good instructor)", show.p = FALSE)
```




```
# Calculate lower and upper bounds for 'aroma' odds
aroma.odds.lower <- exp(aroma.logodds.lower)
aroma.odds.upper <- exp(aroma.logodds.upper)

# Display the confidence interval
paste("(", aroma.odds.lower, ",", aroma.odds.upper, ")")
```

```
[1] "( 1216.28279431965 , 14610.9170234023 )"
```


```
# Plot odds of being a good instructor
plot_model(model1, show.values = TRUE,
           title = "Odds (Good instructor)", show.p = FALSE)
```



```
# Add predicted probabilities
data_aroma_after <- data_aroma %>%
  mutate(logodds.Good = predict(model1, type = "response")) %>%
  mutate(odds.Good = exp(logodds.Good)) %>%
  mutate(probs.Good = fitted(model1))

# Plot the relationship between 'aroma' and probability of being a good instructor
ggplot(data = data_aroma_after, aes(x = aroma, y = probs.Good)) +
```

```
geom_smooth(method = "glm",
            method.args = list(family = "binomial"),
            se = FALSE) +
labs(x = "aroma", y = "Probability of instructor being Good")
```




### 3.4 Flavor and Qualityclass

```
# Select 'flavor' and 'Qualityclass' columns
data_flavor <- data %>%
  dplyr::select(flavor, Qualityclass)

# Create a boxplot of 'flavor' across different 'Qualityclass' levels
p2 <- ggplot(data = data_flavor, aes(x = Qualityclass, y = flavor, fill = Qualityclass)) +
  geom_boxplot() +
  labs(x = "Qualityclass", y = "flavor") +
  theme(legend.position = "none")
p2
```





```
# Fit logistic regression model with 'flavor' predictor and 'Qualityclass' response
model2 <- glm(Qualityclass ~ flavor, data = data_flavor,
              family = binomial(link = "logit"))
model2 %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ flavor, family = binomial(link = "logit"),
    data = data_flavor)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-80.7455	6.1070	-13.22	<2e-16 ***
flavor	10.7238	0.8097	13.24	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.53 on 725 degrees of freedom  
 Residual deviance: 563.98 on 724 degrees of freedom

AIC: 567.98


Number of Fisher Scoring iterations: 6

```
# Calculate lower and upper bounds for 'flavor' log-odds
mod2.coef.logodds <- model2 %>%
  summary() %>%
  coef()
flavor.logodds.lower <- mod2.coef.logodds["flavor", "Estimate"] -
  1.96 * mod2.coef.logodds["flavor", "Std. Error"]
flavor.logodds.upper <- mod2.coef.logodds["flavor", "Estimate"] +
  1.96 * mod2.coef.logodds["flavor", "Std. Error"]

# Display the confidence interval
paste("(", flavor.logodds.lower, ",", flavor.logodds.upper, ")")
```

```
[1] "( 9.1369267163387 , 12.3107616668265 )"
```

```
# Plot log-odds of being a good instructor
plot_model(model2, show.values = TRUE, transform = NULL,
  title = "Log-Odds (Good instructor)", show.p = FALSE)
```




```
# Calculate lower and upper bounds for 'flavor' odds
flavor.odds.lower <- exp(flavor.logodds.lower)
flavor.odds.upper <- exp(flavor.logodds.upper)

# Display the confidence interval
paste("(", flavor.odds.lower, ",", flavor.odds.upper, ")")
```

```
[1] "( 9292.16374923337 , 222073.051342309 )"
```


```
# Plot odds of being a good instructor
plot_model(model2, show.values = TRUE,
           title = "Odds (Good instructor)", show.p = FALSE)
```



```
# Add predicted probabilities
data_flavor_after <- data_flavor %>%
  mutate(logodds.Good = predict(model2, type = "response")) %>%
  mutate(odds.Good = exp(logodds.Good)) %>%
  mutate(probs.Good = fitted(model2))

# Plot the relationship between 'flavor' and probability of being a good instructor
ggplot(data = data_flavor_after, aes(x = flavor, y = probs.Good)) +
```


```
geom_smooth(method="glm",
            method.args = list(family="binomial"),
            se = FALSE) +
labs(x = "flavor", y = "Probability of instructor being Good")
```



### 3.5 Acidity and Qualityclass

```
# Select 'acidity' and 'Qualityclass' columns
data_acidity <- data %>%
  dplyr::select(acidity, Qualityclass)

# Create a boxplot of 'acidity' across different 'Qualityclass' levels
p3 <- ggplot(data = data_acidity, aes(x = Qualityclass, y = acidity, fill = Qualityclass))
  geom_boxplot() +
  labs(x = "Qualityclass", y = "acidity")+
  theme(legend.position = "none")
p3
```



```
# Fit logistic regression model with 'acidity' predictor and 'Qualityclass' response
model3 <- glm(Qualityclass ~ acidity, data = data_acidity,
              family = binomial(link = "logit"))
model3 %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ acidity, family = binomial(link = "logit"),
    data = data_acidity)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-58.4727	4.4807	-13.05	<2e-16 ***
acidity	7.7736	0.5945	13.08	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.5 on 725 degrees of freedom  
 Residual deviance: 675.4 on 724 degrees of freedom

AIC: 679.4


Number of Fisher Scoring iterations: 5

```
# Calculate lower and upper bounds for 'acidity' log-odds
mod3.coef.logodds <- model3 %>%
  summary() %>%
  coef()
acidity.logodds.lower <- mod3.coef.logodds["acidity", "Estimate"] -
  1.96 * mod3.coef.logodds["acidity", "Std. Error"]
acidity.logodds.upper <- mod3.coef.logodds["acidity", "Estimate"] +
  1.96 * mod3.coef.logodds["acidity", "Std. Error"]

# Display the confidence interval
paste("(", acidity.logodds.lower, ",", acidity.logodds.upper, ")")
```

```
[1] "( 6.60847256990056 , 8.93873878119764 )"
```

```
# Plot log-odds of being a good instructor
plot_model(model3, show.values = TRUE, transform = NULL,
  title = "Log-Odds (Good instructor)", show.p = FALSE)
```




```
# Calculate lower and upper bounds for 'acidity' odds
acidity.odds.lower <- exp(acidity.logodds.lower)
acidity.odds.upper <- exp(acidity.logodds.upper)

# Display the confidence interval
paste("(", acidity.odds.lower, ",", acidity.odds.upper, ")")
```

```
[1] "( 741.349793486999 , 7621.57851325419 )"
```


```
# Plot odds of being a good instructor
plot_model(model3, show.values = TRUE,
           title = "Odds (Good instructor)", show.p = FALSE)
```



```
# Add predicted probabilities
data_acidity_after <- data_acidity %>%
  mutate(logodds.Good = predict(model3, type = "response")) %>%
  mutate(odds.Good = exp(logodds.Good)) %>%
  mutate(probs.Good = fitted(model3))

# Plot the relationship between 'acidity' and probability of being a good instructor
ggplot(data = data_acidity_after, aes(x = acidity, y = probs.Good)) +
```

```
geom_smooth(method = "glm",
            method.args = list(family = "binomial"),
            se = FALSE) +
labs(x = "acidity", y = "Probability of instructor being Good")
```




### 3.6 Category 2 type defects and Qualityclass

```
# Select 'category_two_defects' and 'Qualityclass' columns
data_defects <- data %>%
  dplyr::select(category_two_defects, Qualityclass)

# Create a boxplot of 'category_two_defects' across different 'Qualityclass' levels
p4 <- ggplot(data = data_defects, aes(x = Qualityclass, y = category_two_defects, fill = Q
  geom_boxplot() +
  labs(x = "Qualityclass", y = "defects")+
  theme(legend.position = "none")
p4
```





```
# Fit logistic regression model with 'category_two_defects' predictor and 'Qualityclass' r
model5 <- glm(Qualityclass ~ category_two_defects, data = data_defects,
              family = binomial(link = "logit"))
model5 %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ category_two_defects, family = binomial(link = "logit"),
     data = data_defects)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	0.32616	0.10558	3.089	0.00201	**
category_two_defects	-0.08010	0.02999	-2.671	0.00757	**

— — —

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 1003.53 on 725 degrees of freedom
Residual deviance: 996.31 on 724 degrees of freedom
```

AIC: 1000.3


Number of Fisher Scoring iterations: 4

```
# Calculate lower and upper bounds for 'category_two_defects' log-odds
mod5.coef.logodds <- model5 %>%
  summary() %>%
  coef()
defects.logodds.lower <- mod5.coef.logodds["category_two_defects", "Estimate"] -
  1.96 * mod5.coef.logodds["category_two_defects", "Std. Error"]
defects.logodds.upper <- mod5.coef.logodds["category_two_defects", "Estimate"] +
  1.96 * mod5.coef.logodds["category_two_defects", "Std. Error"]

# Display the confidence interval
paste("(", defects.logodds.lower, ",", defects.logodds.upper, ")")
```

```
[1] "( -0.138879080560977 , -0.0213191899815584 )"
```

```
# Plot log-odds of being a good instructor
plot_model(model5, show.values = TRUE, transform = NULL,
  title = "Log-Odds (Good instructor)", show.p = FALSE)
```



```
# Calculate lower and upper bounds for 'category_two_defects' odds
exp(mod5.coef.logodds)
```


	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	1.3856408	1.111352	21.96245288	1.002008
category_two_defects	0.9230248	1.030444	0.06919116	1.007594

```
defects.odds.lower <- exp(defects.logodds.lower)
defects.odds.upper <- exp(defects.logodds.upper)
```

```
# Display the confidence interval
paste("(", defects.odds.lower, ",", defects.odds.upper, ")")
```


```
[1] "( 0.870333262305556 , 0.978906457563423 )"
```

```
# Plot odds of being a good instructor
plot_model(model5, show.values = TRUE,
           title = "Odds (Good instructor)", show.p = FALSE)
```



```
# Add predicted probabilities
data_defects_after <- data_defects %>%
  mutate(logodds.Good = predict(model5, type = "response")) %>%
  mutate(odds.Good = exp(logodds.Good)) %>%
  mutate(probs.Good = fitted(model5))

# Plot the relationship between 'category_two_defects' and probability of being a good ins
ggplot(data = data_defects_after, aes(x = category_two_defects, y = probs.Good)) +
  geom_smooth(method="glm",
             method.args = list(family="binomial"),
             se = FALSE) +
  labs(x = "defects", y = "Probability of instructor being Good")
```




### 3.7 Altitude mean meters and Qualityclass

```
# Select 'altitude_mean_meters' and 'Qualityclass' columns
data_altitude <- data %>%
  dplyr::select(altitude_mean_meters, Qualityclass)

# Create a boxplot of 'altitude_mean_meters' across different 'Qualityclass' levels
p5 <- ggplot(data = data_altitude, aes(x = Qualityclass, y = altitude_mean_meters, fill =
```

```
geom_boxplot() +
labs(x = "Qualityclass", y = "altitude")+
theme(legend.position = "none")
p5
```



```
# Fit logistic regression model with 'altitude_mean_meters' predictor and 'Qualityclass' r
model4 <- glm(Qualityclass ~ altitude_mean_meters, data = data_altitude,
              family = binomial(link = "logit"))
model4 %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ altitude_mean_meters, family = binomial(link = "logit"),
    data = data_altitude)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.1292	0.0755	1.711	0.087 .
altitude_mean_meters	0.3531	0.0774	4.562	5.06e-06 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.53 on 725 degrees of freedom  
Residual deviance: 981.86 on 724 degrees of freedom  
AIC: 985.86


Number of Fisher Scoring iterations: 4

```
# Calculate lower and upper bounds for 'altitude_mean_meters' log-odds
mod4.coef.logodds <- model4 %>%
  summary() %>%
  coef()
altitude.logodds.lower <- mod4.coef.logodds["altitude_mean_meters", "Estimate"] -
  1.96 * mod4.coef.logodds["altitude_mean_meters", "Std. Error"]
altitude.logodds.upper <- mod4.coef.logodds["altitude_mean_meters", "Estimate"] +
  1.96 * mod4.coef.logodds["altitude_mean_meters", "Std. Error"]

# Display the confidence interval
paste("(", altitude.logodds.lower, ",", altitude.logodds.upper, ")")
```

```
[1] "( 0.201429364953632 , 0.504856686546667 )"
```

```
# Plot log-odds of being a good instructor
plot_model(model4, show.values = TRUE, transform = NULL,
  title = "Log-Odds (Good instructor)", show.p = FALSE)
```



```
# Calculate lower and upper bounds for 'altitude_mean_meters' odds
exp(mod4.coef.logodds)
```


	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	1.137935	1.078426	5.536717	1.090904
altitude_mean_meters	1.423535	1.080480	95.801741	1.000005

```
altitude.odds.lower <- exp(altitude.logodds.lower)
altitude.odds.upper <- exp(altitude.logodds.upper)

# Display the confidence interval
paste("(", altitude.odds.lower, ", ", altitude.odds.upper, ")")
```

```
[1] "( 1.22314983676597 , 1.65674806915918 )"
```


```
# Plot odds of being a good instructor
plot_model(model4, show.values = TRUE,
            title = "Odds (Good instructor)", show.p = FALSE)
```



```
# Add predicted probabilities
data_altitude_after <- data_altitude %>%
  mutate(logodds.Good = predict(model4), type = "response") %>%
  mutate(odds.Good = exp(logodds.Good)) %>%
  mutate(probs.Good = fitted(model4))

# Plot the relationship between 'altitude_mean_meters' and probability of being a good ins
ggplot(data = data_altitude_after, aes(x = altitude_mean_meters, y = probs.Good)) +
  geom_smooth(method="glm",
             method.args = list(family="binomial"),
             se = FALSE) +
  labs(x = "altitude", y = "Probability of instructor being Good")
```






### 3.8 Harvested and Qualityclass

```
# Select 'harvested' and 'Qualityclass' columns and generate a contingency table.
data_harvested <- data %>%
  dplyr::select(harvested, Qualityclass)
data_harvested %>%
  tabyl(harvested, Qualityclass) %>%
  adorn_percentages() %>%
  adorn_pct_formatting() %>%
  adorn_ns()
```

harvested	Poor		Good	
2010	28.6%	(4)	71.4	(10)
2011	32.0%	(8)	68.0	(17)
2012	55.4%	(98)	44.6	(79)
2013	50.9%	(57)	49.1	(55)
2014	47.5%	(77)	52.5	(85)
2015	40.6%	(39)	59.4	(57)
2016	35.0%	(28)	65.0	(52)
2017	52.1%	(25)	47.9	(23)
2018	33.3%	(4)	66.7%	(8)

```
# Create a barplot of 'harvested' across different 'Qualityclass' levels
p6 <- ggplot(data_harvested, aes(x = Qualityclass, y = after_stat(prop), group = harvested))
  geom_bar(position = "dodge", stat = "count") +
  labs(y = "Proportion")
p6
```



```
# Fit logistic regression model with 'harvested' predictor and 'Qualityclass' response
model_harvested <- glm(Qualityclass ~ harvested, data = data_harvested,
  family = binomial(link = "logit"))
model_harvested %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ harvested, family = binomial(link = "logit"),
  data = data_harvested)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.9163	0.5916	1.549	0.1214
harvested2011	-0.1625	0.7306	-0.222	0.8240
harvested2012	-1.1318	0.6106	-1.854	0.0638 .

harvested2013	-0.9520	0.6211	-1.533	0.1253
harvested2014	-0.8174	0.6122	-1.335	0.1818
harvested2015	-0.5368	0.6270	-0.856	0.3920
harvested2016	-0.2973	0.6364	-0.467	0.6404
harvested2017	-0.9997	0.6584	-1.518	0.1289
harvested2018	-0.2231	0.8515	-0.262	0.7933

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.53 on 725 degrees of freedom  
 Residual deviance: 985.86 on 717 degrees of freedom  
 AIC: 1003.9

Number of Fisher Scoring iterations: 4

```
# Extract coefficients from the model and calculate their confidence intervals.
model_harvested_coef_logodds <- model_harvested %>%
  summary() %>%
  coef()
model_harvested_coef_logodds
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.9162907	0.5916076	1.5488150	0.12142620
harvested2011	-0.1625189	0.7306319	-0.2224361	0.82397439
harvested2012	-1.1318104	0.6106242	-1.8535303	0.06380639
harvested2013	-0.9520088	0.6210678	-1.5328581	0.12531083
harvested2014	-0.8174449	0.6121693	-1.3353248	0.18177005
harvested2015	-0.5368011	0.6270442	-0.8560818	0.39195256
harvested2016	-0.2972515	0.6363526	-0.4671177	0.64041570
harvested2017	-0.9996723	0.6583903	-1.5183582	0.12892412
harvested2018	-0.2231436	0.8514690	-0.2620689	0.79326831

```
confint_logodds <- confint(model_harvested)
confint_logodds
```

	2.5 %	97.5 %
(Intercept)	-0.1788402	2.209833043
harvested2011	-1.6742313	1.245165574

```

harvested2012 -2.4551656 0.003857604
harvested2013 -2.2918766 0.205965175
harvested2014 -2.1431438 0.321691403
harvested2015 -1.8858074 0.634595032
harvested2016 -1.6606350 0.894930001
harvested2017 -2.4001925 0.236998318
harvested2018 -1.9324563 1.477106332

```

```

# Plot log-odds of being a good instructor
plot_model(model_harvested, show.values = TRUE, transform = NULL,
           title = "Log-Odds (Good instructor)", show.p = FALSE)

```



```

# Transform the coefficients into odds ratios and obtain their confidence intervals
model_harvested_coef_odds <- model_harvested %>%
  summary() %>%
  coef() %>%
  exp()
model_harvested_coef_odds

```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	2.5000000	1.806891	4.7058905	1.129106

```

harvested2011 0.8500000 2.076392 0.8005661 2.279542
harvested2012 0.3224490 1.841580 0.1566831 1.065886
harvested2013 0.3859649 1.860914 0.2159177 1.133501
harvested2014 0.4415584 1.844428 0.2630727 1.199338
harvested2015 0.5846154 1.872069 0.4248234 1.479867
harvested2016 0.7428571 1.889576 0.6268063 1.897269
harvested2017 0.3680000 1.931680 0.2190713 1.137604
harvested2018 0.8000000 2.343086 0.7694580 2.210610

```

```
exp(confint_logodds)
```

```


                2.5 %   97.5 %
(Intercept) 0.83623950 9.114195
harvested2011 0.18745223 3.473510
harvested2012 0.08584898 1.003865
harvested2013 0.10107661 1.228710
harvested2014 0.11728554 1.379459
harvested2015 0.15170652 1.886258
harvested2016 0.19001827 2.447164
harvested2017 0.09070049 1.267439
harvested2018 0.14479211 4.380252

```

```


# Plot odds of being a good instructor
plot_model(model_harvested, show.values = TRUE,
           title = "Odds (Good instructor)", show.p = FALSE)

```




```
# Calculate log odds, odds, and probabilities and store them
data_harvested_after <- data_harvested %>%
  mutate(logodds.Good = predict(model_harvested, type = "response")) %>%
  mutate(odds.Good = exp(logodds.Good)) %>%
  mutate(probs.Good = fitted(model_harvested))

# Generate a predictive plot
plot_model(model_harvested, type = "pred",
  terms = c("harvested"))
```




### 3.9 Plot Arrange

```
# Arrange multiple plots  
grid.arrange(p1, p2, p3, p4, p5, ncol = 3)
```



```
grid.arrange(p0, p6)
```





## 4 Formal Analysis

### 4.1 Principal Component Analysis

Based on the correlation matrix, it is evident that some predictors exhibit high correlation. Therefore, we adopt principal component analysis (PCA) to help address multicollinearity, thereby enhancing the stability and interpretability of the model.

```
# Principal principal component analysis (PCA) for 'aroma', 'flavor' and 'acidity'
data_pca <- data %>%
  dplyr::select(aroma, flavor, acidity, Qualityclass)
data_scaled <- scale(data_pca[, -4])
pca_result <- prcomp(data_scaled)
summary(pca_result)
```

Importance of components:

	PC1	PC2	PC3
Standard deviation	1.5170	0.6790	0.48747
Proportion of Variance	0.7671	0.1537	0.07921
Cumulative Proportion	0.7671	0.9208	1.00000

The cumulative proportion of the three predictor variables adds up to 1, indicating that these three principal components fully explain the variability in the original data without losing information. Therefore, adopting principal component analysis is justified.

```
# Predict PCA components and choose the first two components
pca_result_selected <- predict(pca_result, newdata = data_scaled)[, 1:2]

# Combine PCA components with other variables
data_pca_final <- data.frame(pca_result_selected, country_category = data$country_category)

# Retrieve column names of the new data frame
names(data_pca_final)
```

```
[1] "PC1"                "PC2"                "country_category"
[4] "category_two_defects" "altitude_mean_meters" "harvested"
[7] "Qualityclass"
```

## 4.2 Model Selection

```
# Conduct an origin model
model_full_after <- glm(Qualityclass ~ country_category + aroma + flavor + acidity + category_two_defects + altitude_mean_meters + harvested,
                        family = binomial(link = "logit"), data = data)

# Summarize the model
model_full_after %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ country_category + aroma + flavor +
     acidity + category_two_defects + altitude_mean_meters + harvested,
     family = binomial(link = "logit"), data = data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-136.36544	11.87262	-11.486	< 2e-16	***
country_categoryColombia	0.19714	3.83740	0.051	0.9590	
country_categoryThailand	0.38580	3.91529	0.099	0.9215	
country_categoryother	-1.98656	3.82843	-0.519	0.6038	
aroma	5.60251	0.89564	6.255	3.97e-10	***
flavor	7.41691	1.02128	7.262	3.80e-13	***
acidity	5.13315	0.84515	6.074	1.25e-09	***
category_two_defects	0.08151	0.05280	1.544	0.1227	
altitude_mean_meters	0.23102	0.14289	1.617	0.1059	
harvested2011	-0.09750	1.08274	-0.090	0.9282	
harvested2012	0.07982	0.91032	0.088	0.9301	
harvested2013	0.22982	0.90818	0.253	0.8002	
harvested2014	0.94042	0.92551	1.016	0.3096	
harvested2015	0.68207	0.93880	0.727	0.4675	
harvested2016	1.76568	0.98341	1.795	0.0726	.
harvested2017	1.23785	0.98308	1.259	0.2080	
harvested2018	2.65646	1.30031	2.043	0.0411	*

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.53 on 725 degrees of freedom  
Residual deviance: 396.36 on 709 degrees of freedom  
AIC: 430.36

Number of Fisher Scoring iterations: 7

```
# Perform stepwise variable selection using AIC
stepAIC(model_full_after)
```

Start: AIC=430.36

```
Qualityclass ~ country_category + aroma + flavor + acidity +
  category_two_defects + altitude_mean_meters + harvested
```

	Df	Deviance	AIC
<none>		396.36	430.36
- category_two_defects	1	398.77	430.77
- altitude_mean_meters	1	398.97	430.97
- harvested	8	419.15	437.15
- country_category	3	435.09	463.09
- acidity	1	438.67	470.67
- aroma	1	443.99	475.99
- flavor	1	462.60	494.60

```
Call: glm(formula = Qualityclass ~ country_category + aroma + flavor +
  acidity + category_two_defects + altitude_mean_meters + harvested,
  family = binomial(link = "logit"), data = data)
```

Coefficients:

(Intercept)	country_categoryColombia	country_categoryThailand
-136.36544	0.19714	0.38580
country_categoryother	aroma	flavor
-1.98656	5.60251	7.41691
acidity	category_two_defects	altitude_mean_meters
5.13315	0.08151	0.23102
harvested2011	harvested2012	harvested2013
-0.09750	0.07982	0.22982
harvested2014	harvested2015	harvested2016
0.94042	0.68207	1.76568
harvested2017	harvested2018	
1.23785	2.65646	

Degrees of Freedom: 725 Total (i.e. Null); 709 Residual

Null Deviance: 1004

Residual Deviance: 396.4 AIC: 430.4

```
# Fit logistic regression model with PCA components
pca_model <- glm(Qualityclass ~ ., data = data_pca_final, family = binomial(link = "logit")

# Summarize the model
pca_model %>%
  summary()
```

Call:

```
glm(formula = Qualityclass ~ ., family = binomial(link = "logit"),
    data = data_pca_final)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	1.03127	3.89993	0.264	0.7914
PC1	0.64526	0.05280	12.222	<2e-16 ***
PC2	-0.01683	0.04837	-0.348	0.7278
country_categoryColombia	0.13759	3.81016	0.036	0.9712
country_categoryThailand	0.27024	3.88866	0.069	0.9446
country_categoryother	-2.13272	3.80058	-0.561	0.5747
category_two_defects	0.08881	0.05245	1.693	0.0904 .
altitude_mean_meters	0.18852	0.14058	1.341	0.1799
harvested2011	-0.16663	1.08668	-0.153	0.8781
harvested2012	0.07030	0.91843	0.077	0.9390
harvested2013	0.12648	0.91488	0.138	0.8900
harvested2014	0.91531	0.93415	0.980	0.3272
harvested2015	0.69761	0.94612	0.737	0.4609
harvested2016	1.75244	0.99220	1.766	0.0774 .
harvested2017	1.28878	0.98794	1.305	0.1921
harvested2018	2.48395	1.29643	1.916	0.0554 .

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1003.53 on 725 degrees of freedom  
 Residual deviance: 399.57 on 710 degrees of freedom  
 AIC: 431.57

Number of Fisher Scoring iterations: 7

```
pca_model_summary <- glance(pca_model)
kable(pca_model_summary, digits = 2)
```

null.deviance	df.null	logLik	AIC	BIC	deviance	df.residual	nobs
1003.53	725	-199.79	431.57	504.97	399.57	710	726

```
# Perform stepwise variable selection using AIC
stepAIC(pca_model)
```

Start: AIC=431.57

Qualityclass ~ PC1 + PC2 + country\_category + category\_two\_defects +  
altitude\_mean\_meters + harvested

	Df	Deviance	AIC
- PC2	1	399.69	429.69
- altitude_mean_meters	1	401.37	431.37
<none>		399.57	431.57
- category_two_defects	1	402.48	432.48
- harvested	8	422.90	438.90
- country_category	3	441.10	467.10
- PC1	1	915.48	945.48

Step: AIC=429.69

Qualityclass ~ PC1 + country\_category + category\_two\_defects +  
altitude\_mean\_meters + harvested

	Df	Deviance	AIC
- altitude_mean_meters	1	401.55	429.55
<none>		399.69	429.69
- category_two_defects	1	402.65	430.65
- harvested	8	423.99	437.99
- country_category	3	441.13	465.13
- PC1	1	916.55	944.55

Step: AIC=429.55

Qualityclass ~ PC1 + country\_category + category\_two\_defects +  
harvested

	Df	Deviance	AIC
<none>		401.55	429.55

```

- category_two_defects 1 404.86 430.86
- harvested 8 424.94 436.94
- country_category 3 449.38 471.38
- PC1 1 931.59 957.59

```

```

Call: glm(formula = Qualityclass ~ PC1 + country_category + category_two_defects +
  harvested, family = binomial(link = "logit"), data = data_pca_final)

```

Coefficients:

```

              (Intercept)              PC1 country_categoryColombia
              1.39357              0.64662              -0.09348
country_categoryThailand country_categoryother category_two_defects
              -0.14086              -2.48718              0.09380
harvested2011 harvested2012 harvested2013
              -0.19451              -0.01981              0.12008
harvested2014 harvested2015 harvested2016
              0.88638              0.65190              1.63107
harvested2017 harvested2018
              1.35077              2.30535

```

Degrees of Freedom: 725 Total (i.e. Null); 712 Residual

Null Deviance: 1004

Residual Deviance: 401.6 AIC: 429.6

After reducing dimensionality using PCA, we selected the model with the lowest AIC, which is considered the optimal model.

```

# Final Logistic Regression Model for Qualityclass Prediction
optimal_model <- glm(Qualityclass ~ PC1 + country_category + category_two_defects + harvested,
  family = binomial(link = "logit"), data = data_pca_final)
optimal_model %>%
  summary()

```

Call:

```

glm(formula = Qualityclass ~ PC1 + country_category + category_two_defects +
  harvested, family = binomial(link = "logit"), data = data_pca_final)

```

Coefficients:

```

              Estimate Std. Error z value Pr(>|z|)
(Intercept)      1.39357      3.96280   0.352   0.7251

```

```

PC1                0.64662    0.05276   12.256   <2e-16 ***
country_categoryColombia -0.09348    3.88231   -0.024    0.9808
country_categoryThailand -0.14086    3.94836   -0.036    0.9715
country_categoryother   -2.48718    3.86848   -0.643    0.5203
category_two_defects    0.09380    0.05198    1.804    0.0712 .
harvested2011          -0.19451    1.07478   -0.181    0.8564
harvested2012          -0.01981    0.91171   -0.022    0.9827
harvested2013           0.12008    0.91107    0.132    0.8951
harvested2014           0.88638    0.93212    0.951    0.3416
harvested2015           0.65190    0.94433    0.690    0.4900
harvested2016           1.63107    0.97215    1.678    0.0934 .
harvested2017           1.35077    0.98144    1.376    0.1687
harvested2018           2.30535    1.27791    1.804    0.0712 .

```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 1003.53  on 725  degrees of freedom
Residual deviance:  401.55  on 712  degrees of freedom
AIC: 429.55

```

Number of Fisher Scoring iterations: 7

```

optimal_model_summary <- glance(optimal_model)
kable(optimal_model_summary,digits =2)

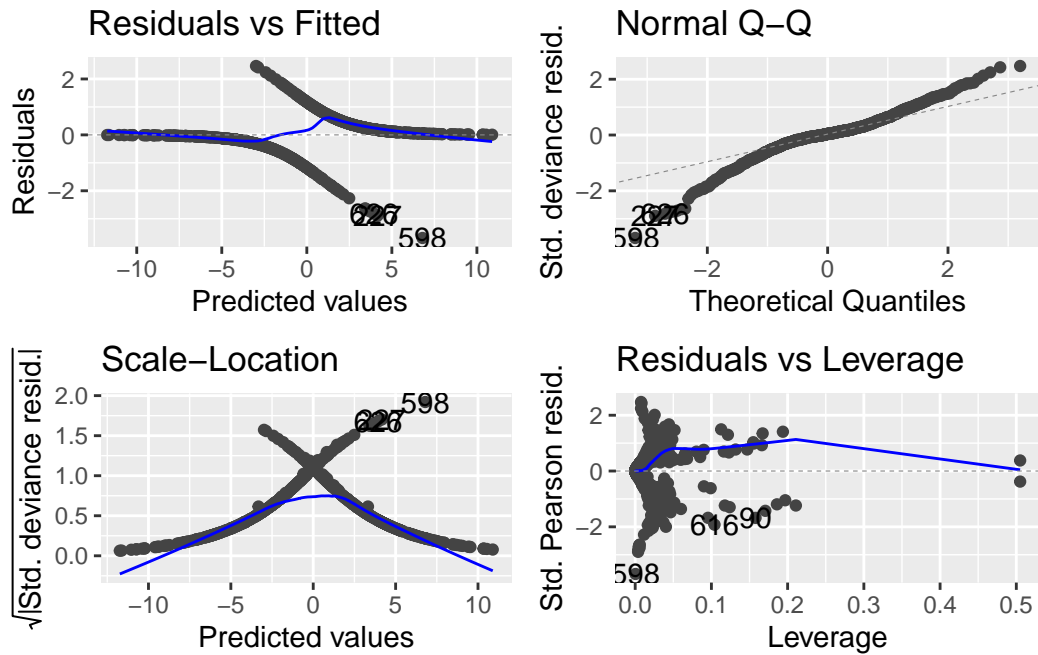
```

null.deviance	df.null	logLik	AIC	BIC	deviance	df.residual	nobs
1003.53	725	-200.78	429.55	493.78	401.55	712	726

```

# Check the assumptions
autoplot(optimal_model)

```



$$\text{Qualityclass} = \beta_0 + \beta_1 \times \text{PC1} + \beta_2 \times \text{country\_category} + \beta_3 \times \text{category\_two\_defects} + \beta_4 \times \text{harvested} + \epsilon$$

- *Qualityclass* is the response variable
- *PC1* is a variable derived from reducing the dimensions of *aroma*, *flavor*, and *acidity*
- *country\_category*, *category\_two\_defects*, and *harvested* are the predictor variables
- $\beta_0$  to  $\beta_4$  are the coefficients of the model
- $\epsilon$  is the error term